
Advanced Prediction Models

Deep Learning, Graphical Models and Reinforcement
Learning

Recap: Why Graphical Models

- We have seen deep learning techniques for unstructured data
 - Predominantly vision and text/audio
 - We will see control in the last part of the course
 - (Reinforcement Learning)
- For structured data, graphical models are the most versatile framework
 - Successfully applications:
 - Kalman filtering in engineering
 - Decoding in cell phones (channel codes)
 - Hidden Markov models for time series
 - Clustering, regression, classification ...

Recap: Graphical Models Landscape

- Three key parts:
 - Representation
 - Capture uncertainty (joint distribution)
 - Capture conditional independences (metadata)
 - Visualization of metadata for a distribution
 - Inference
 - Efficient methods for computing marginal or conditional distributions quickly
 - Learning
 - Learning the parameters of the distribution can deal with prior knowledge and missing data

Today's Outline

- Inference
 - Factor Graph
 - Variable Elimination
- Inference using Belief Propagation
- Inference using Markov Chain Monte Carlo

Inference

Based on notes from Bjoern Andres and Bernt Schiele (2016)

Inference Objectives

- Let $\bar{X} = X_1, \dots, X_D$ be a random vector.
- Let $\bar{X} \in \mathfrak{X}$ and $X_i \in \mathfrak{X}_i$
- Given $P(\bar{X})$ compute functions of it

Inference Objectives

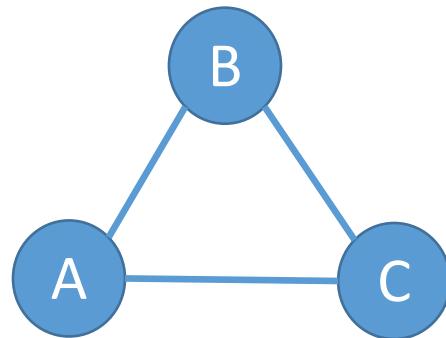
- Let $\bar{X} = X_1, \dots, X_D$ be a random vector.
- Let $\bar{X} \in \mathfrak{X}$ and $X_i \in \mathfrak{X}_i$
- Given $P(\bar{X})$ compute functions of it
 - Example, find
 - Mode $\bar{x}^* \in \operatorname{argmax}_{\bar{x} \in \mathfrak{X}} P(\bar{x})$
 - Mean $E[g(\bar{x})] = \sum_{\bar{x} \in \mathfrak{X}} g(\bar{x})P(\bar{x})$
 - A marginal $\operatorname{argmax}_{x_i \in \mathfrak{X}_i} \sum_{x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_D} P(\bar{x})$
 - A conditional $P(X_i | x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_D)$

Algorithms for Inference

- Variable Elimination
- Belief Propagation
- Sampling based methods (MCMC)

Factor Graphs

- For both DPGM and UPGMs, factorization is simply not specified by the graph!
- Consider the following example graph

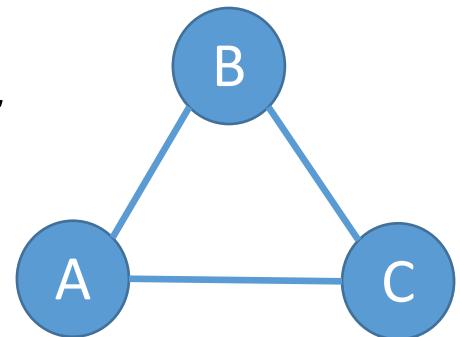


- It could be $P(a, b, c) = \frac{1}{Z} \phi(a, b, c)$
- Or it could be $P(a, b, c) = \frac{1}{Z} \phi_1(a, b) \phi_2(b, c) \phi_3(c, a)$

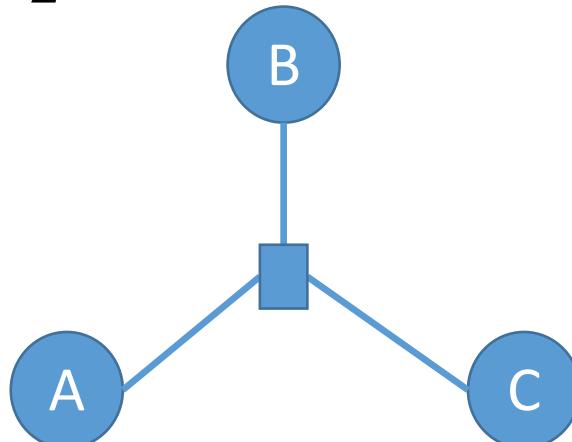
Factor Graph for UPGM

- Hence, we define new graphs called **factor graphs**

- Consider a square node for each factor



- Then, $P(a, b, c) = \frac{1}{Z} \phi(a, b, c)$ can be represented by

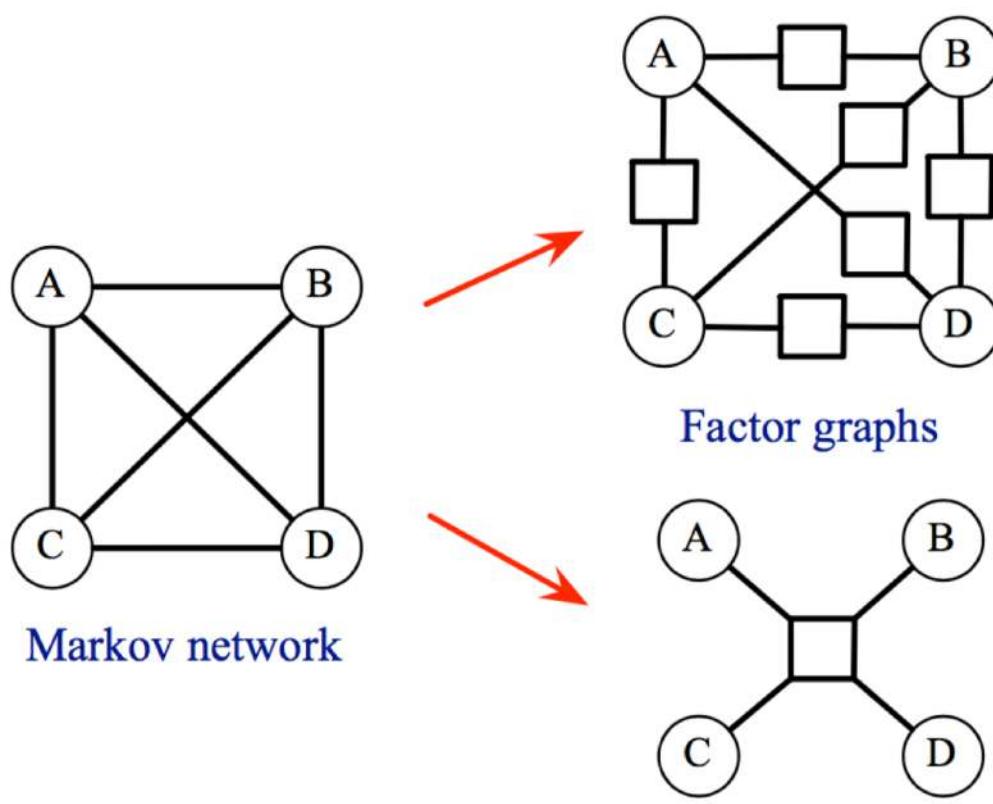


Factor Graph

- factor graphs capture the factorization in the graph itself
- For a function $f(x_1, \dots, x_D) = \prod_i \phi_i(\mathcal{X}_i)$ the factor graph has a square node for each factor $\phi_i(\mathcal{X}_i)$ and a circular variable node for each variable x_j
- Factor graphs will allow us to define inference algorithms for both DPGMs and UPGMs
 - Just a more richer way of drawing graphs for $P(X)$

Factor Graphs for a UPGM

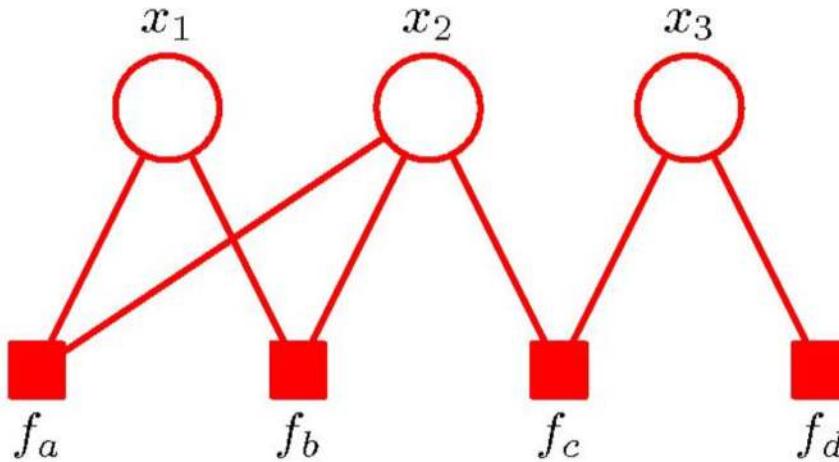
- The following example shows two factor graphs for the same UPGM



Factor Graph Example (I)

- Which distribution does the following graph correspond to?

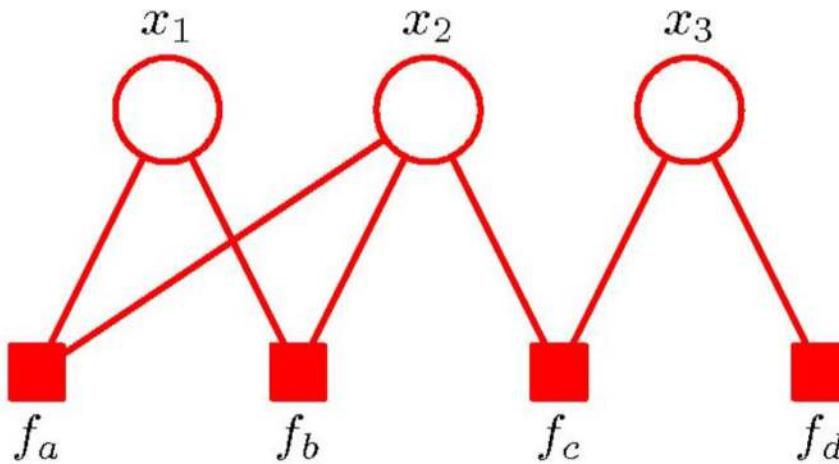
We will use f or ϕ to denote factors



We will use lower case to minimize notation clutter

Factor Graph Example (I)

- Which distribution does the following graph correspond to?



- It corresponds to
 - $P(x_1, x_2, x_3) = \frac{1}{Z} f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3)$

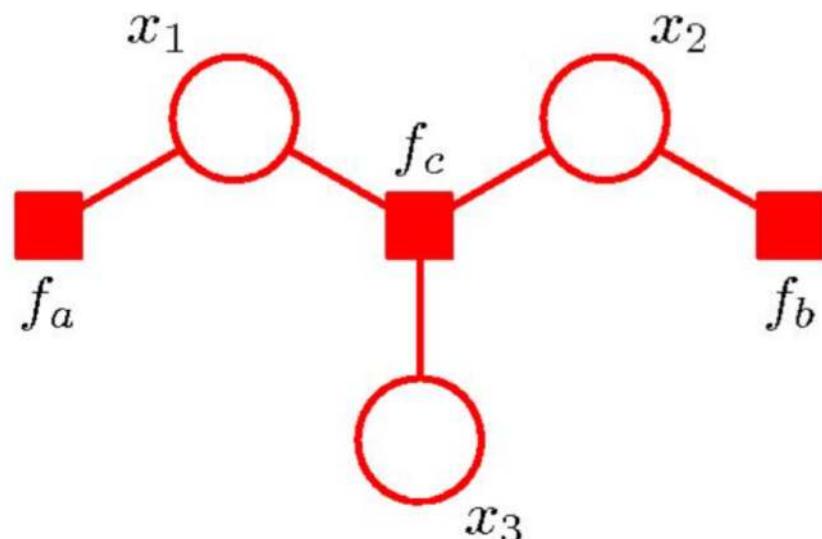
Factor Graph Example (II)

- What is the factor graph for the distribution
 - $P(x_1, x_2, x_3) = \frac{1}{Z} f_c(x_3 | x_1, x_2) f_a(x_1) f_b(x_2)$

Factor Graph Example (II)

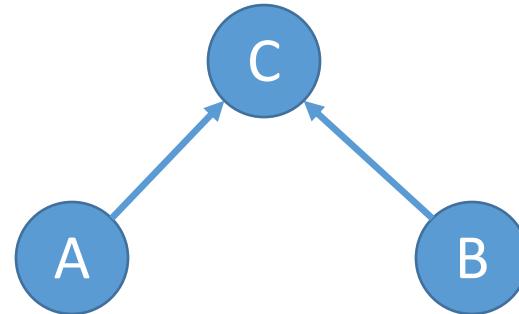
- What is the factor graph for the distribution
 - $P(x_1, x_2, x_3) = \frac{1}{Z} f_c(x_3 | x_1, x_2) f_a(x_1) f_b(x_2)$
- The following is the desired factor graph

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

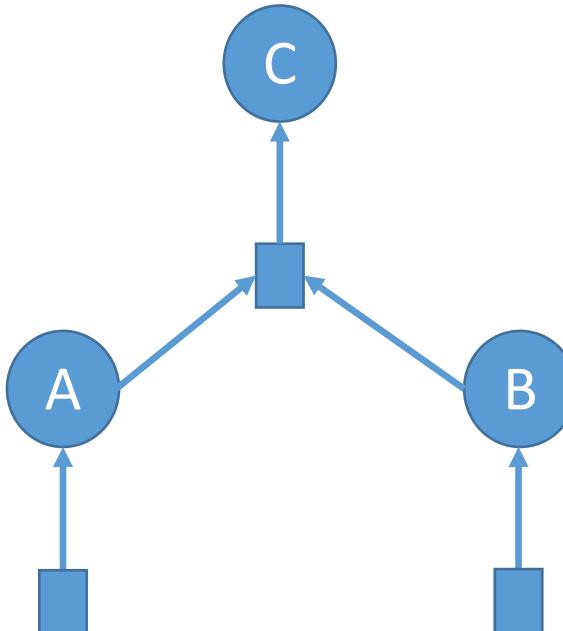


Factor Graph for DPGM

- We can do this for DPGMs as well (although redundant)



- Consider the graph on the right



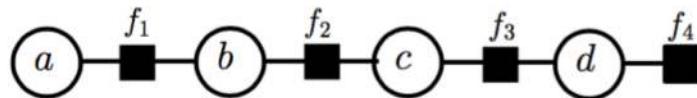
- Its factor graph representation is shown below

Inference using Variable Elimination

- It is a very simple idea, which is
 - Don't sum over all configurations simultaneously
 - Do it one variable at a time
- Works for DPGMs and UPGMs

Variable Elimination Example

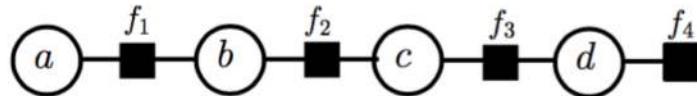
We will use lower case to minimize notation clutter



This can be for a DPGM or a UPGM

Variable Elimination Example

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)



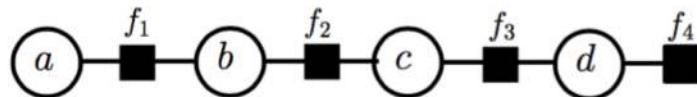
This can be for a DPGM or a UPGM

$$p(a, b, c, d) = \frac{1}{Z} f_1(a, b) f_2(b, c) f_3(c, d) f_4(d)$$

Objective: Find $p(a, b)$

Variable Elimination Example

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)



$$p(a, b, c, d) = \frac{1}{Z} f_1(a, b) f_2(b, c) f_3(c, d) f_4(d)$$

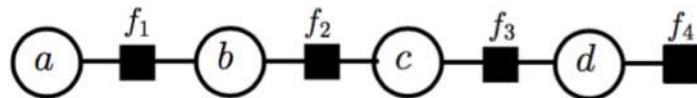
Objective: Find $p(a, b)$

$$\begin{aligned} p(a, b, c) &= \sum_d p(a, b, c, d) \\ &= \frac{1}{Z} f_1(a, b) f_2(b, c) \underbrace{\sum_d f_3(c, d) f_4(d)}_{\mu_{d \rightarrow c}(c)} \end{aligned}$$

(compute this for all c)

Variable Elimination Example

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)



$$p(a, b, c, d) = \frac{1}{Z} f_1(a, b) f_2(b, c) f_3(c, d) f_4(d)$$

Objective: Find $p(a, b)$

$$\begin{aligned} p(a, b, c) &= \sum_d p(a, b, c, d) \\ &= \frac{1}{Z} f_1(a, b) f_2(b, c) \underbrace{\sum_d f_3(c, d) f_4(d)}_{\mu_{d \rightarrow c}(c)} \quad (\text{compute this for all } c) \end{aligned}$$

$$p(a, b) = \sum_c p(a, b, c) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_c f_2(b, c) \mu_{d \rightarrow c}(c)}_{\mu_{c \rightarrow b}(b)} \quad (\text{compute this for all } b)$$

Questions?

Today's Outline

- Inference
 - Factor Graph
 - Variable Elimination
- Inference using Belief Propagation
- Inference using Markov Chain Monte Carlo

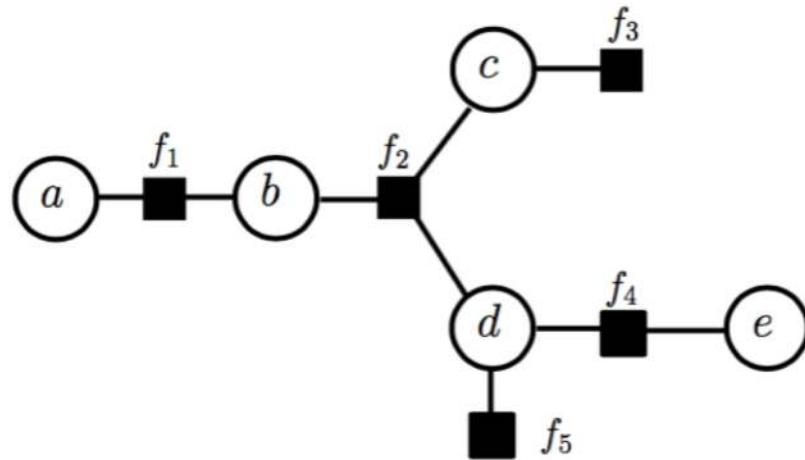
Inference using Belief Propagation

Belief Propagation (BP)

- Generalizes the idea of Variable Elimination
 - Also called the Sum-Product Algorithm
-
- Will give exact answers (marginals, conditionals) on factor graphs that are trees
 - Can also be used for general graphs but may give wrong answers

BP Example: Compute a Marginal

consider a branching graph:



with factors

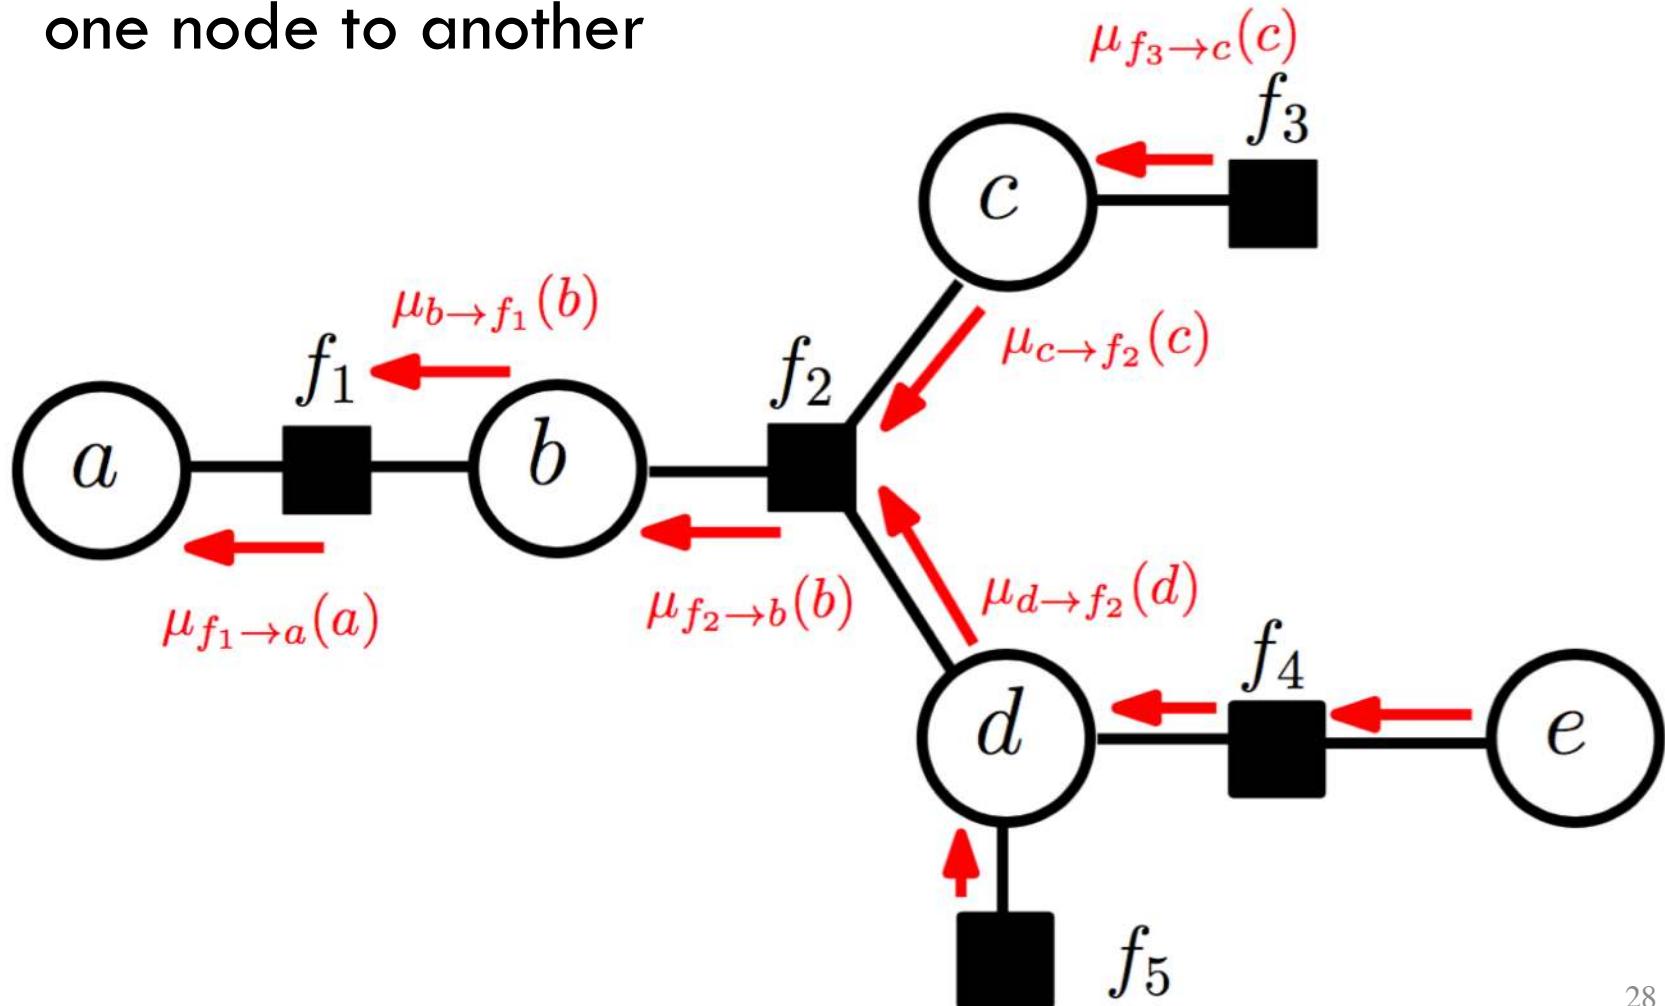
$$f_1(a, b) f_2(b, c, d) f_3(c) f_4(d, e) f_5(d)$$

For example: find marginal $p(a, b)$

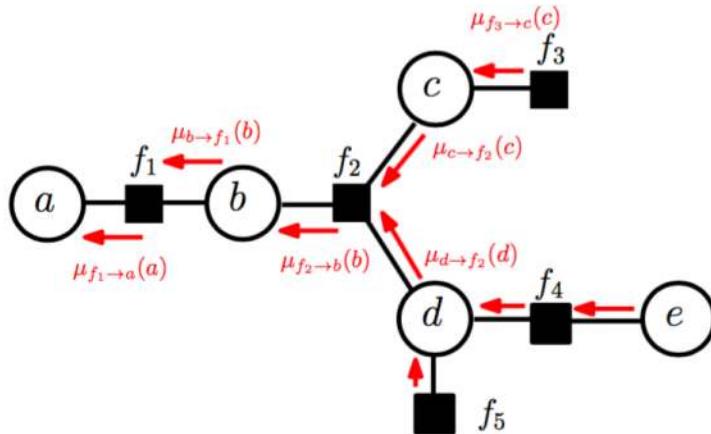
- We will introduce the notion of
 - messages, and
 - message passing

BP Example: Messages

- Messages are functions (vectors) that are passed from one node to another



BP Example: Messages



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c,d,e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f2 \rightarrow b}(b)}$$

$$\mu_{f2 \rightarrow b}(b) = \sum_{c,d} f_2(b, c, d) \underbrace{f_3(c)}_{\mu_{c \rightarrow f2}(c)} \underbrace{f_5(d) \sum_e f_4(d, e)}_{\mu_{d \rightarrow f2}(d)}$$

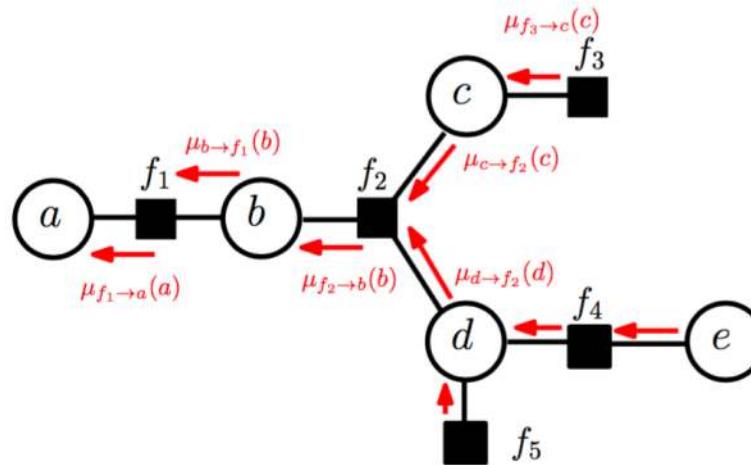
BP Example: Message from Factor to Variable

Here (repeated from last slide):

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c,d} f_2(b, c, d) \underbrace{f_3(c)}_{\mu_{c \rightarrow f_2}(c)} \underbrace{f_5(d)}_{\mu_{d \rightarrow f_2}(d)} \sum_e f_4(d, e)$$

$$\mu_{\textcolor{red}{f_2}} \rightarrow b(b) = \sum_{c,d} \textcolor{red}{f_2}(b, c, d) \mu_{c \rightarrow \textcolor{red}{f_2}}(c) \mu_{d \rightarrow \textcolor{red}{f_2}}(d)$$

BP Example: Message from Factor to Variable



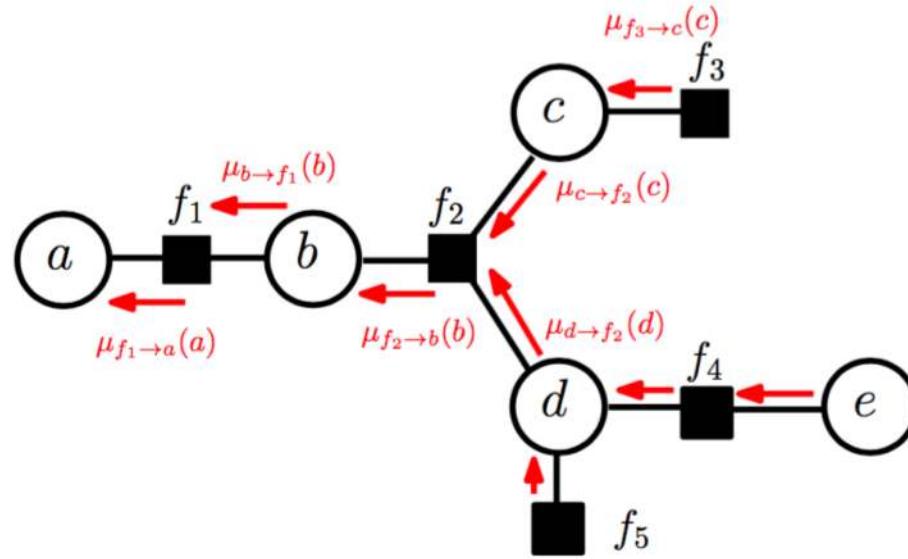
Here (repeated from last slide):

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c,d} f_2(b, c, d) \mu_{c \rightarrow f_2}(c) \mu_{d \rightarrow f_2}(d)$$

more general:

$$\mu_{\mathbf{f} \rightarrow x}(x) = \sum_{y \in \mathcal{X}_{\mathbf{f}} \setminus x} \phi_{\mathbf{f}}(\mathcal{X}_{\mathbf{f}}) \prod_{y \in \{\text{ne}(\mathbf{f}) \setminus x\}} \mu_{y \rightarrow \mathbf{f}}(y)$$

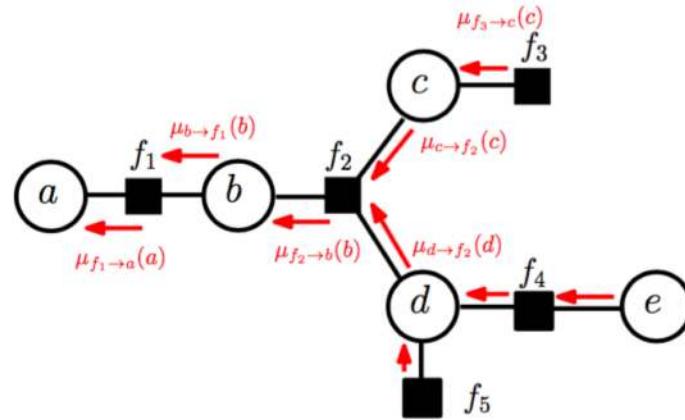
BP Example: Message from Variable to Factor



$$\mu_{d \rightarrow f_2}(d) = \underbrace{f_5(d)}_{\mu_{f_5 \rightarrow d}(d)} \underbrace{\sum_e f_4(d, e)}_{\mu_{f_4 \rightarrow d}(d)}$$

$$\mu_{\textcolor{red}{d} \rightarrow f_2}(\textcolor{red}{d}) = \mu_{f_5 \rightarrow \textcolor{red}{d}}(\textcolor{red}{d}) \mu_{f_4 \rightarrow \textcolor{red}{d}}(\textcolor{red}{d})$$

BP Example: Message from Variable to Factor



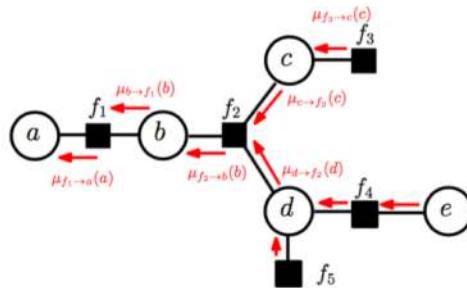
Here (repeated from last slide):

$$\mu_{d \rightarrow f_2}(d) = \mu_{f_5 \rightarrow d}(d) \mu_{f_4 \rightarrow d}(d)$$

General:

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{\text{ne}(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$

BP Example: Compute a Different Marginal



If we want to compute the marginal $p(a)$
(use factor-to-variable message):

$$p(a) = \frac{1}{Z} \mu_{f_1 \rightarrow a}(a) = \underbrace{\sum_b f_1(a, b) \mu_{b \rightarrow f_1}(b)}_{\mu_{f_1 \rightarrow a}(a)} \frac{1}{Z}$$

which we could also view as

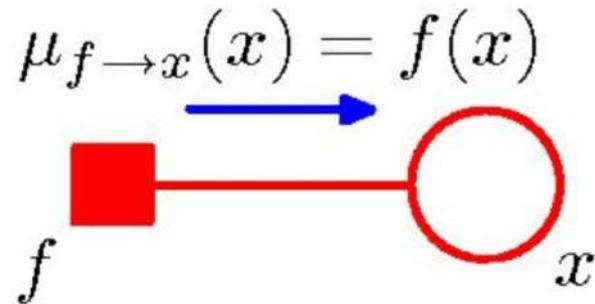
$$p(a) = \frac{1}{Z} \sum_b f_1(a, b) \underbrace{\mu_{b \rightarrow f_1}(b)}_{\mu_{f_2 \rightarrow b}(b)}$$

Belief Propagation Algorithm

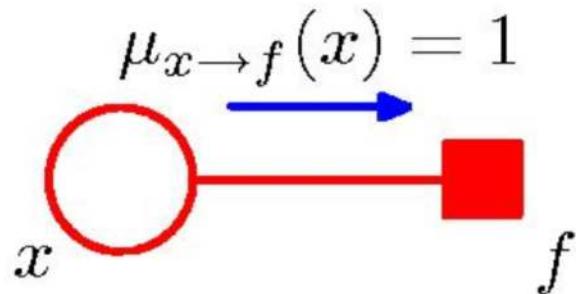
- We described the concept of ‘messages’ via an example (computing marginals for a given factor graph)
- Now we will summarize the algorithm in general
- It has three key ingredients
 - Initialization
 - Variable to factor message
 - Factor to variable message
- Don’t forget the original objective: efficient inference

BP: Initialization

- Messages from extremal/leaf node factors are initialized to be the factor itself

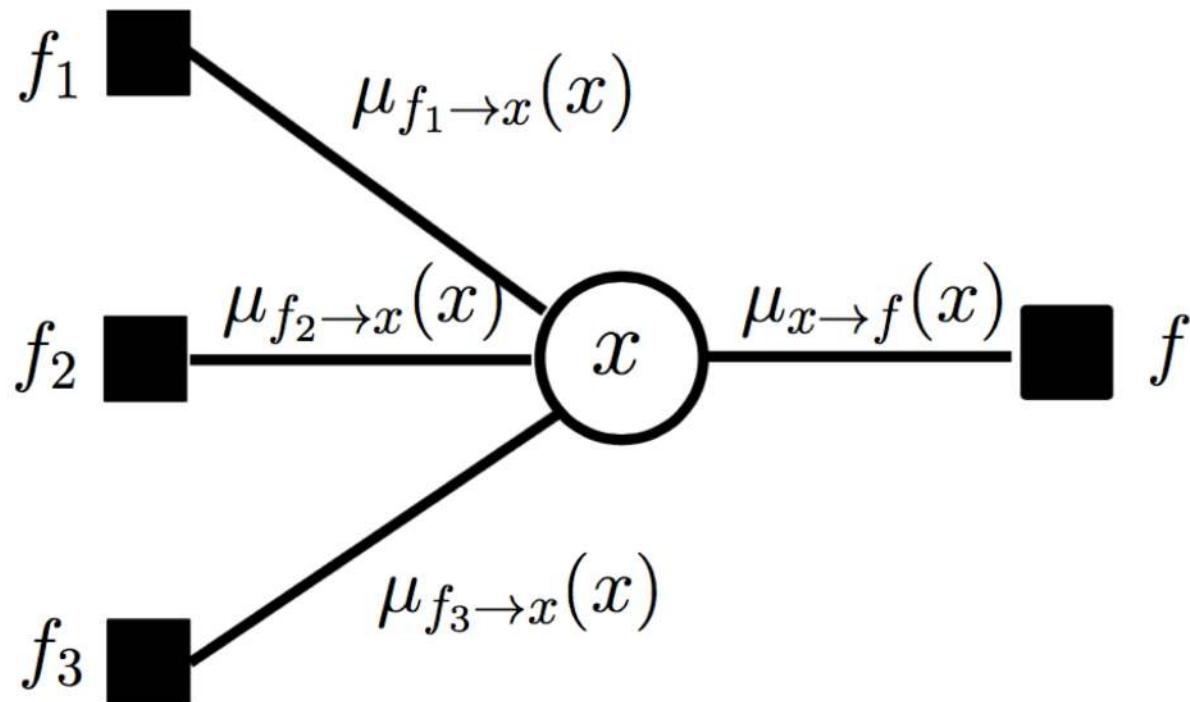


- Messages from extremal/leaf node variables are initialized to value 1



BP: Variable to Factor Message

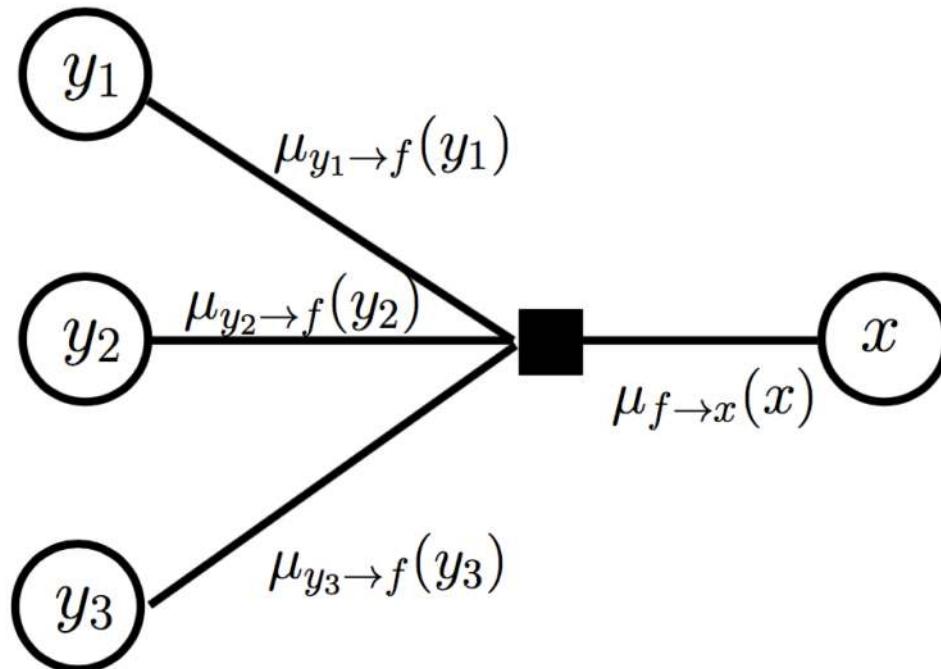
$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{\text{ne}(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$



BP: Factor to Variable Message

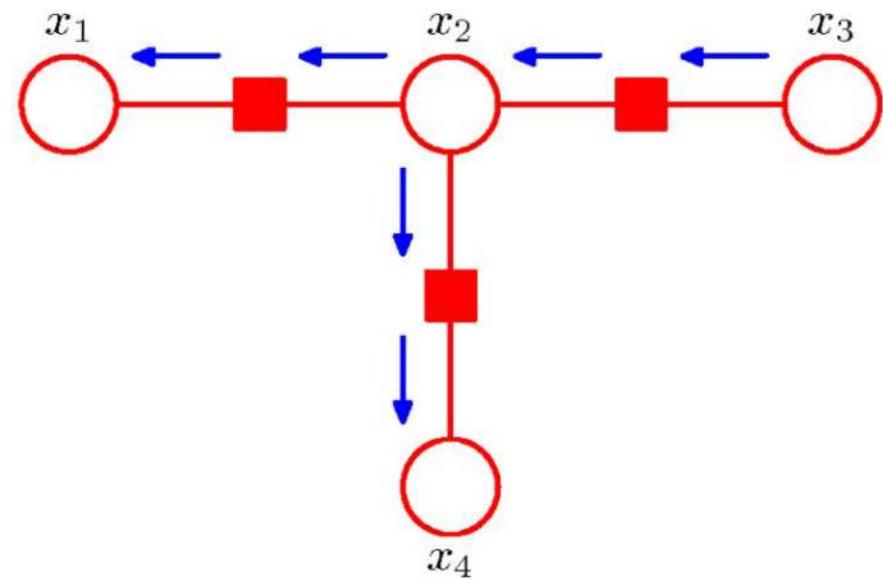
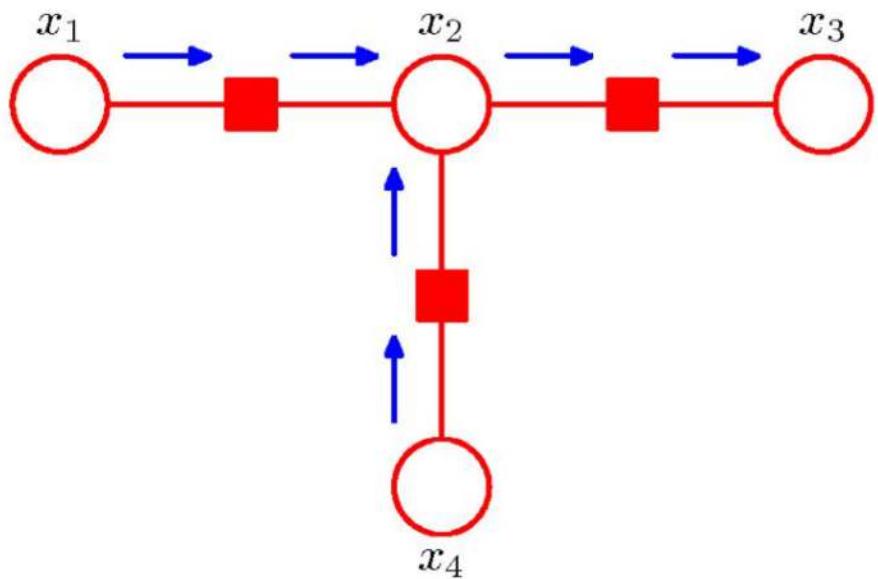
- We sum over all values possible in the scope of the factor

$$\mu_{f \rightarrow x}(x) = \sum_{y \in \mathcal{X}_f \setminus x} \phi_f(\mathcal{X}_f) \prod_{y \in \{\text{ne}(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$



BP: Ordering of Messages

- Messages depend on all incoming messages
- To compute all messages
 - Go from leaves to a designated root (say x_3)
 - Go from the designated root back to leaves

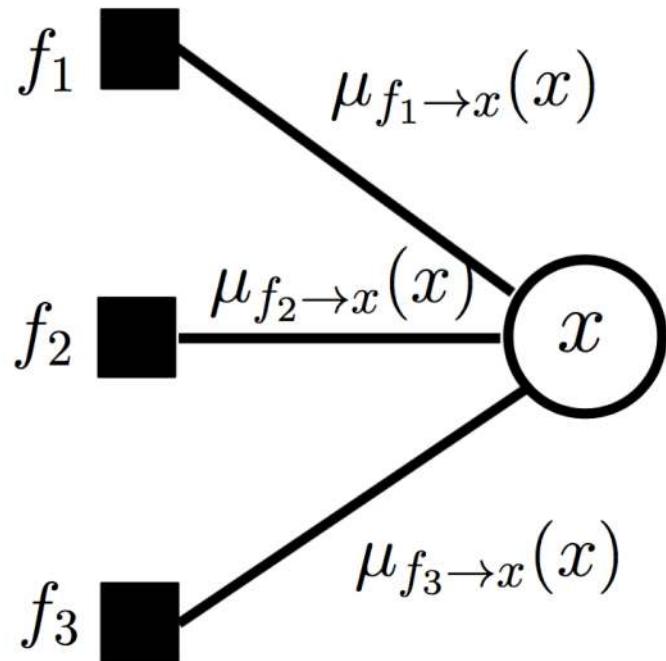


Designated root: x_3

BP: Computing a Marginal

- Marginal is simply the product of messages the variable of interest receives

$$p(x) \propto \prod_{f \in \text{ne}(x)} \mu_{f \rightarrow x}(x)$$



BP: General Factor Graphs

- Is in-exact
- Since it is not clear whether BP is a clear winner for inference with general graphs (among competing algorithms), we will not explore this further.
- See https://en.wikipedia.org/wiki/Belief_propagation for more details

Questions?

Today's Outline

- Inference
 - Factor Graphs
 - Variable Elimination
- Inference using Belief Propagation
- Inference using Markov Chain Monte Carlo

Inference using Markov Chain Monte Carlo

See https://en.wikipedia.org/wiki/Markov_chain_Monte_Carlo

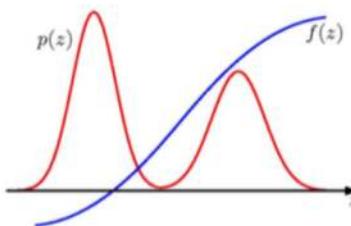
Approximate Inference

- BP and Variable Elimination are exact algorithms
- They work for tree structured factor graphs
- We will resort to numerical sampling to perform approximate inference for general graphical models
 - Essentially, use random sampling to approximate

Sampling

- Many methods in the literature
- Monte Carlo methods
 - MC Averaging and Importance sampling
 - Rejection sampling
- Markov Chain Monte Carlo methods
 - Gibbs sampling
 - Metropolis-Hastings sampling
- ...

Monte Carlo Averaging



We want to evaluate

$$\mathbb{E}[f] = \int f(x)p(x)dx \quad \text{or} \quad \mathbb{E}[f] = \sum_{x \in \mathcal{X}} f(x)p(x)$$

Sampling idea:

- ▶ draw L independent samples x^1, x^2, \dots, x^L from $p(\cdot)$: $x^l \sim p(\cdot)$
- ▶ replace the integral/sum with the finite set of samples

$$\hat{f} = \frac{1}{L} \sum_{l=1}^L f(x^l)$$

- ▶ as long as $x^l \sim p(\cdot)$ then

$$\mathbb{E}[\hat{f}] = \mathbb{E}[f]$$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Monte_Carlo_method

Importance Sampling

- Is a variance reduction technique for MC averaging

use a proposal distribution $q(z)$ from which it is easy to draw samples
express expectation in the form of a finite sum over samples $\{z^l\}$
drawn from $q(z)$:

$$\begin{aligned}\mathbb{E}[f] &= \int f(z)p(z)dz = \int f(z)\frac{p(z)}{q(z)}q(z)dz \\ &\simeq \frac{1}{L} \sum_{l=1}^L \frac{p(z^l)}{q(z^l)} f(z^l)\end{aligned}$$

with importance weights: $r^l = \frac{p(z^l)}{q(z^l)}$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Importance_sampling

Importance Sampling

- If we can only evaluate up to a normalizing constant, then additional tricks needed.

$p(z)$ can be only evaluated up to a normalization constant (unkown):

$$p(z) = \tilde{p}(z)/Z_p$$

$q(z)$ can be also treated in a similar way:

$$q(z) = \tilde{q}(z)/Z_q$$

then:

$$\begin{aligned}\mathbb{E}[f] &= \int f(z)p(z)dz = \frac{Z_q}{Z_p} \int f(z) \frac{\tilde{p}(z)}{\tilde{q}(z)} q(z) dz \\ &\simeq \frac{Z_q}{Z_p} \frac{1}{L} \sum_{l=1}^L \tilde{r}^l f(z^l)\end{aligned}$$

with: $\tilde{r}^l = \frac{\tilde{p}(z^l)}{\tilde{q}(z^l)}$

For example,

$$\frac{Z_p}{Z_q} \simeq \frac{1}{L} \sum_{l=1}^L \tilde{r}^l$$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Importance_sampling

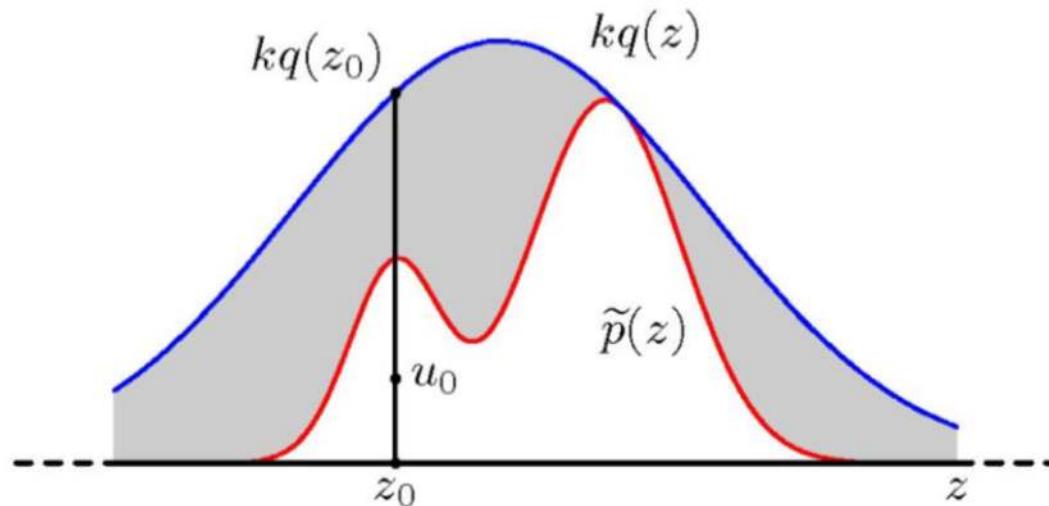
Rejection Sampling

Sample two random variables:

1. $z_0 \sim q(x)$
2. $u_0 \sim [0, kq(z_0)]$ uniform

$q(x)$ is a proposal distribution such that $kq(x) \geq p(x) \forall x$

reject sample z_0 if $u_0 > \tilde{p}(z_0)$



¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Rejection_sampling

Rejection Sampling

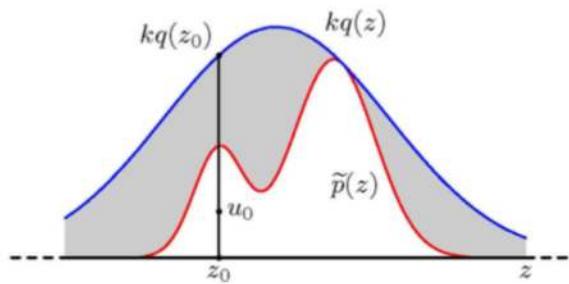
Sample z drawn from q and accepted with probability $\tilde{p}(z)/kq(z)$

So (overall) acceptance probability

$$p(\text{accept}) = \int \frac{\tilde{p}(z)}{kq(z)} q(z) dz = \frac{1}{k} \int \tilde{p}(z) dz$$

So the lower k the better (more acceptance)

- ▶ subject to constraint $kq(z) \geq \tilde{p}(z)$

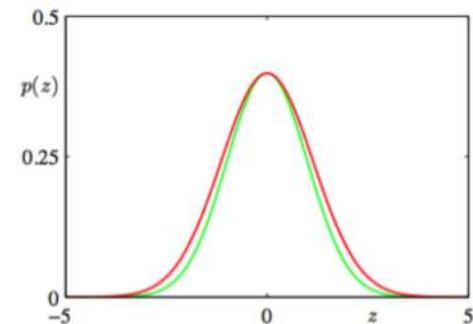


- Impractical in high dimensions (lots of samples will get rejected)

Rejection Sampling

Example:

- ▶ assume $p(x)$ is Gaussian with covariance matrix: $\sigma_p^2 I$
- ▶ assume $q(x)$ is Gaussian with covariance matrix: $\sigma_q^2 I$
- ▶ clearly: $\sigma_q^2 \geq \sigma_p^2$
- ▶ in D dimensions: $k = \left(\frac{\sigma_q}{\sigma_p}\right)^D$

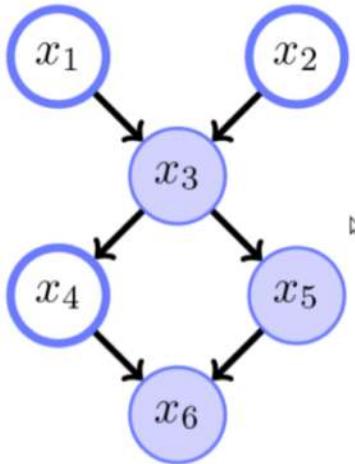


assume:

- ▶ σ_q is 1% larger than σ_p , $D = 1000$
- ▶ then $k = 1.01^{1000} \geq 20000$
- ▶ and $p(\text{accept}) \leq \frac{1}{20000}$

therefore: often impractical to find good proposal distribution $q(x)$ for high dimensions

Gibbs Sampling: Markov Blanket



Sample from this distribution $p(x)$

Idea: Sample sequence x^0, x^1, x^2, \dots by updating one variable at a time

Eg. update x_4 by conditioning on the set of shaded variables **Markov blanket**

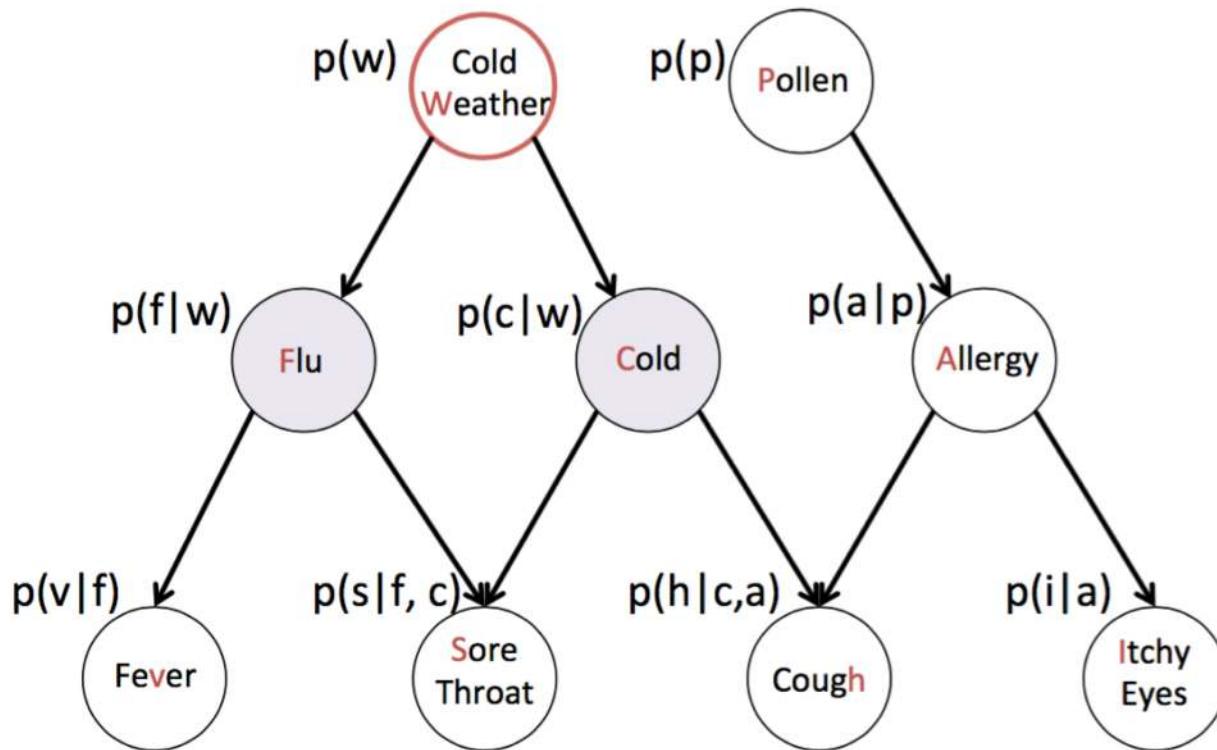
$$p(x_4 | x_1, x_2, x_3, x_5, x_6) = p(x_4 | x_3, x_5, x_6)$$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Gibbs_sampling

Gibbs Sampling Example I

How do we sample a new value for W?

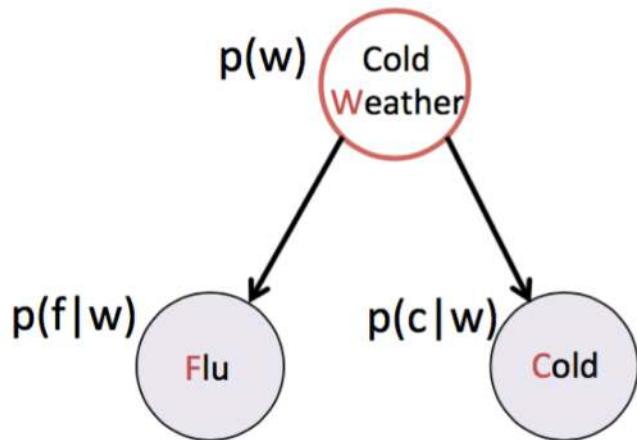


$$P(W=w|F=1, P=1, C=0, \dots, I=0)$$

$$= P(W=w|F=1, C=0)$$

Markov Blanket!

Gibbs Sampling Example I



w	$p(w)$
0	0.4
1	0.6

w	f	$p(f w)$
0	0	0.95
0	1	0.05
1	0	0.80
1	1	0.20

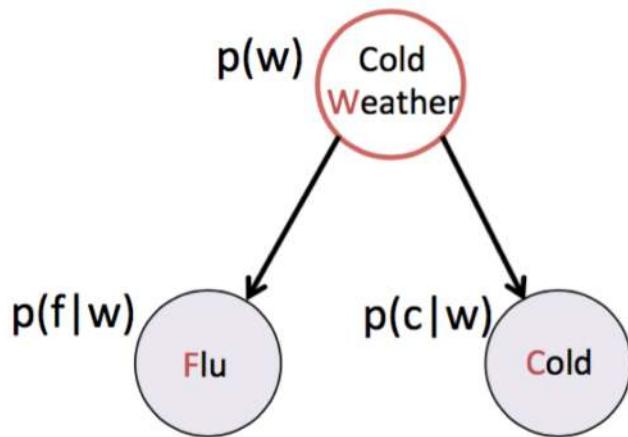
w	c	$p(c w)$
0	0	0.88
0	1	0.12
1	0	0.70
1	1	0.30

$$P(W=w|F=1, P=1, C=0, \dots, I=0)$$

$$= P(W=w|F=1, C=0)$$

$$\propto P(F=1|W=w) * P(C=0|W=w) * P(W=w)$$

Gibbs Sampling Example I



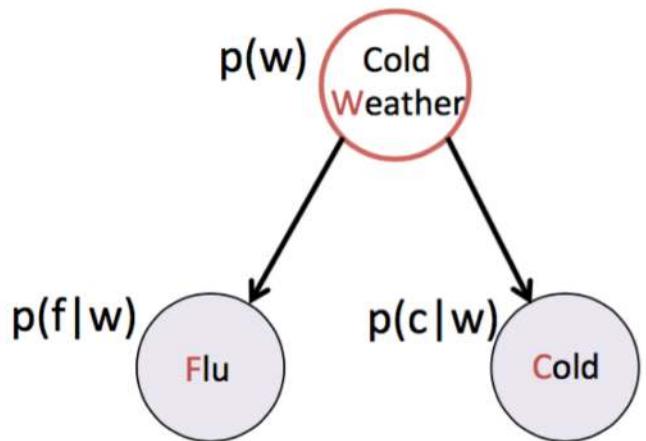
w	p(w)
0	0.40
1	0.60

w	f	p(f w)
0	0	0.95
0	1	0.05
1	0	0.80
1	1	0.20

w	c	p(f w)
0	0	0.88
0	1	0.12
1	0	0.70
1	1	0.30

$$\begin{aligned} & P(W=w | F=1, P=1, C=0, \dots, I=0) \\ &= P(W=w | F=1, C=0) \\ &\propto P(F=1 | W=w) * P(C=0 | W=w) * P(W=w) \\ &= \begin{cases} 0.05 * 0.88 * 0.40, & W = 0 \\ & \vdots \end{cases} \end{aligned}$$

Gibbs Sampling Example I



w	p(w)
0	0.40
1	0.60

w	f	p(f w)
0	0	0.95
0	1	0.05
1	0	0.80
1	1	0.20

w	c	p(f w)
0	0	0.88
0	1	0.12
1	0	0.70
1	1	0.30

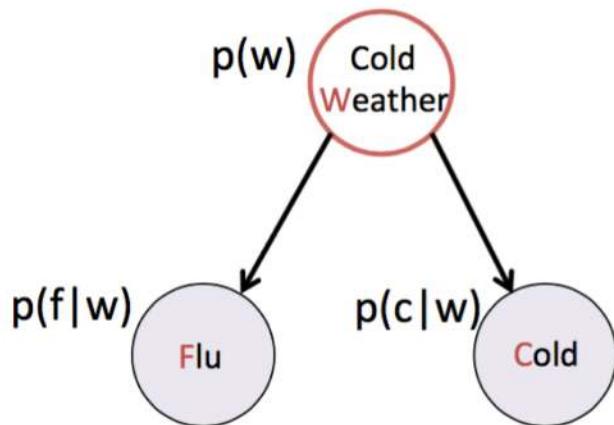
$$P(W=w|F=1, P=1, C=0, \dots, I=0)$$

$$= P(W=w|F=1, C=0)$$

$$\propto P(F=1|W=w)*P(C=0|W=w)*P(W=w)$$

$$= \begin{cases} 0.05 * 0.88 * 0.40, & W = 0 \\ 0.20 * 0.70 * 0.60, & W = 1 \end{cases}$$

Gibbs Sampling Example I



w	p(w)
0	0.40
1	0.60

w	f	p(f w)
0	0	0.95
0	1	0.05
1	0	0.80
1	1	0.20

w	c	p(f w)
0	0	0.88
0	1	0.12
1	0	0.70
1	1	0.30

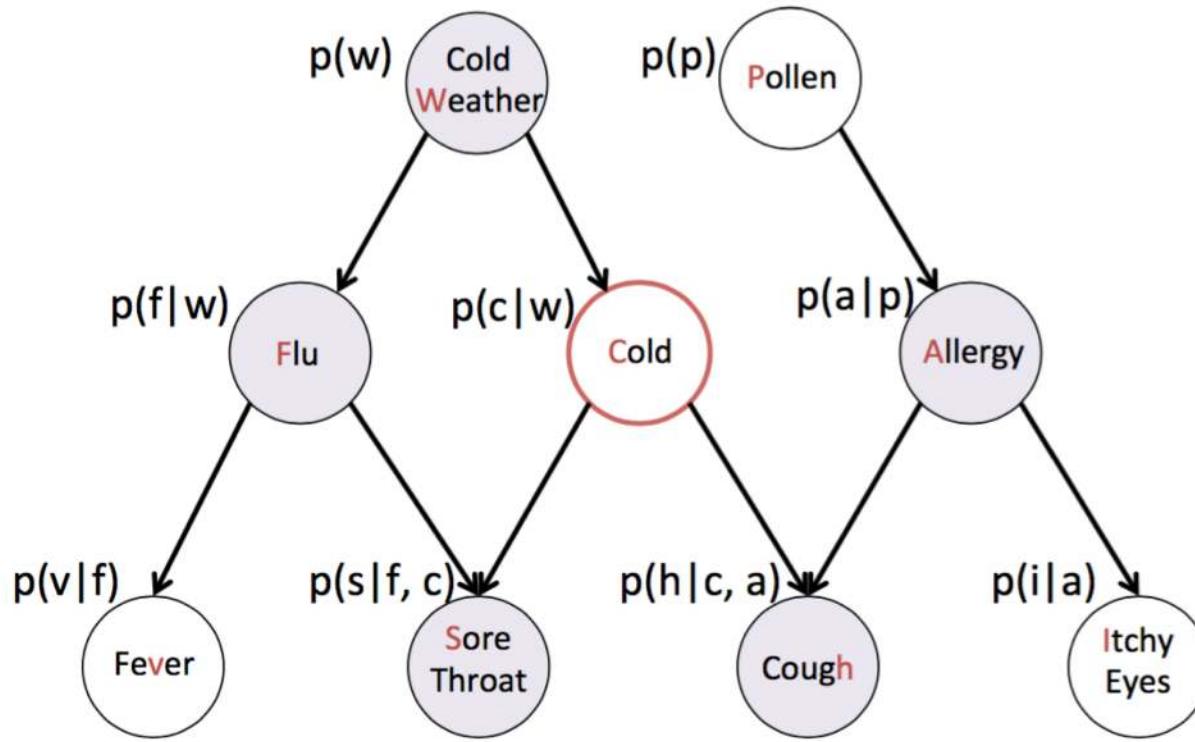
$$P(W=w | F=1, C=0)$$

$$= P(W=w | F=1, C=0)$$
$$\propto P(F=1 | W=w) * P(C=0 | W=w) * P(W=w)$$
$$= \begin{cases} 0.05 * 0.88 * 0.40, & w = 0 \\ 0.20 * 0.70 * 0.60, & w = 1 \end{cases}$$
$$= \begin{cases} 0.0176 / (0.0176 + 0.084), & w = 0 \\ 0.084 / (0.0176 + 0.084), & w = 1 \end{cases}$$
$$= \begin{cases} 0.173, & w = 0 \\ 0.827, & w = 1 \end{cases}$$

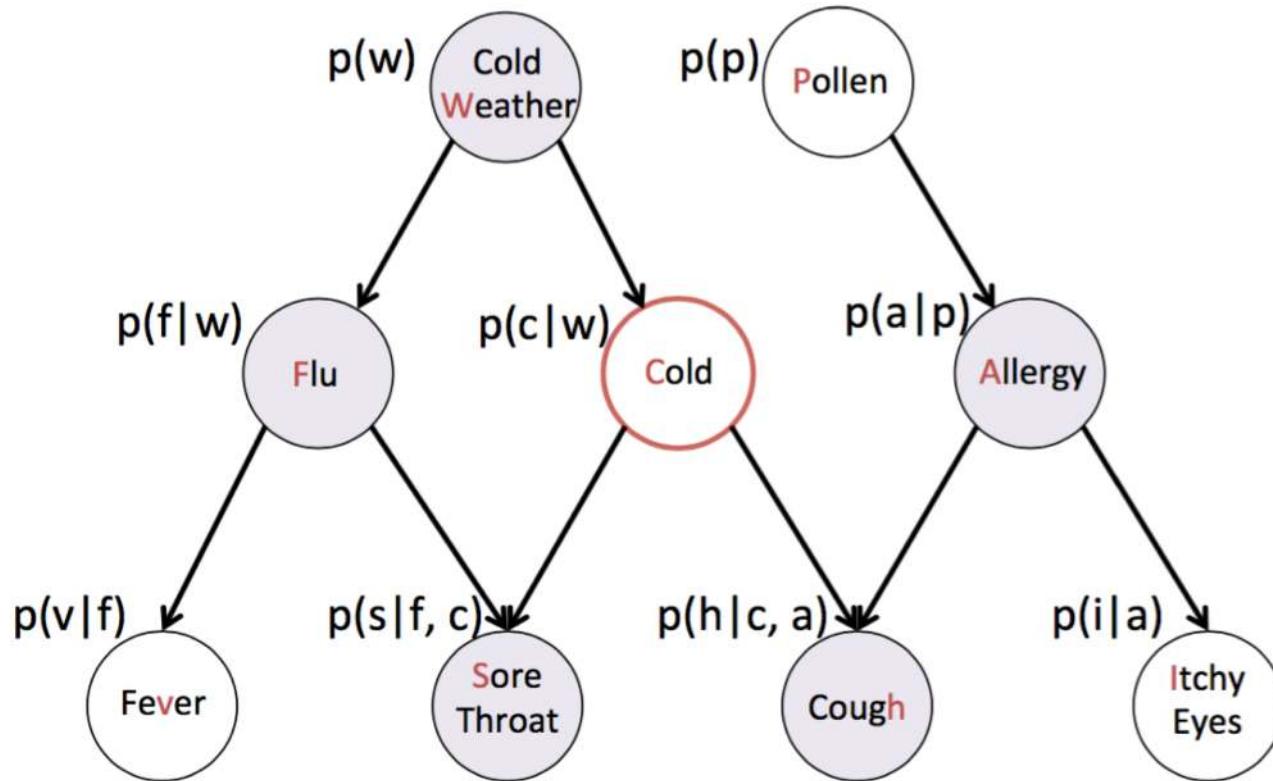
Sample a new w!

Gibbs Sampling Example II

How do we sample a new value for C?



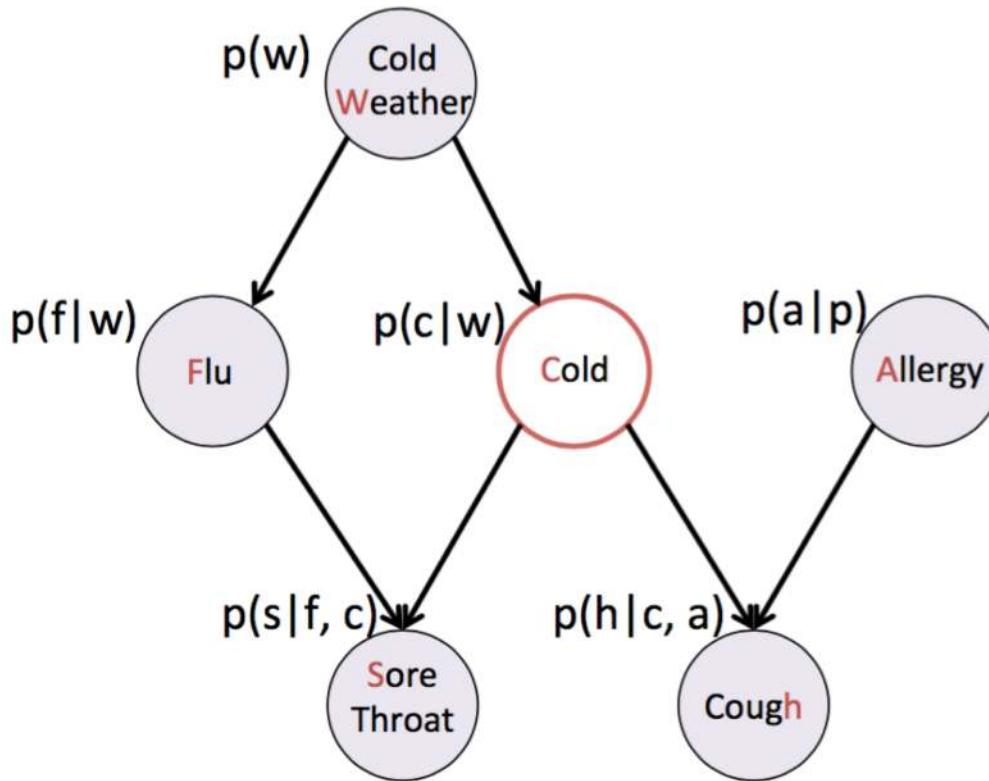
Gibbs Sampling Example II



$$P(C=c \mid W=1, F=1, P=1, \dots, I=0)$$

$$= P(C=c \mid W=1, F=1, S=0, H=1, A=1) \quad \text{Markov Blanket!}$$

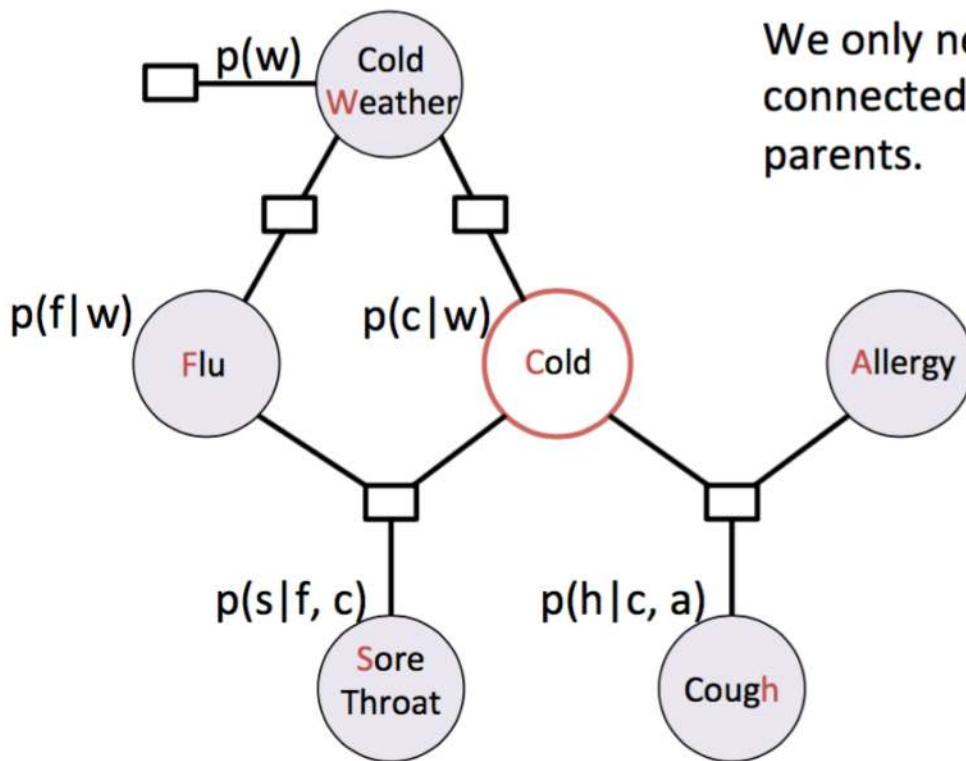
Gibbs Sampling Example II



$$P(C=c \mid W=1, F=1, P=1, \dots, I=0)$$

$$= P(C=c \mid W=1, F=1, S=0, H=1, A=1)$$

Gibbs Sampling Example II



We only need look at the factors connected to the children and the parents.

$$P(C=c | W=1, F=1, P=1, \dots, I=0)$$

$$= P(C=c | W=1, F=1, S=0, H=1, A=1)$$

$$= p(w) p(f | w) p(c | w) p(s | f, c) p(h | c, a)$$

Gibbs Sampling: Conditional Probability

Update x_i

$$p(x_i | x_{\setminus i}) = \frac{1}{Z} p(x_i | pa(x_i)) \prod_{j \in \text{ch}(i)} p(x_j | \text{pa}(x_j))$$

and the normalisation constant is

$$Z = \sum_{x_i} p(x_i | pa(x_i)) \prod_{j \in \text{ch}(i)} p(x_j | \text{pa}(x_j))$$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Gibbs_sampling

Understanding MCMC via Markov Chain Terminology

Sample from a multi-variate distribution

$$p(x) = \frac{1}{Z} p^*(x)$$

with Z intractable to calculate

Idea: Sample from some $q(\mathbf{x} \rightarrow \mathbf{x}')$ with a **stationary distribution**

$$\pi(\mathbf{x}') = \sum_{\mathbf{x}} \pi(\mathbf{x}) q(\mathbf{x} \rightarrow \mathbf{x}') \quad \text{for all } \mathbf{x}'$$

Given $p(x)$ find $q(\mathbf{x} \rightarrow \mathbf{x}')$ such that $\pi(\mathbf{x}) = p(x)$

Gibbs sampling is one instance (that is why it is working)

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Metropolis%E2%80%93Hastings_algorithm

Understanding MCMC via Markov Chain Terminology

Transition probability $q(\mathbf{x} \rightarrow \mathbf{x}')$

Occupancy probability $\pi_t(\mathbf{x})$ at time t

Equilibrium condition on π_t defines stationary distribution $\pi(\mathbf{x})$

Note: stationary distribution depends on choice of $q(\mathbf{x} \rightarrow \mathbf{x}')$

Pairwise detailed balance on states guarantees equilibrium

Gibbs sampling transition probability:

sample each variable given current values of all others

⇒ detailed balance with the true posterior

For Bayesian networks, Gibbs sampling reduces to sampling conditioned on each variable's Markov blanket

Stationary Distribution of a MC

$\pi_t(\mathbf{x})$ = probability in state \mathbf{x} at time t

$\pi_{t+1}(\mathbf{x}')$ = probability in state \mathbf{x}' at time $t + 1$

π_{t+1} in terms of π_t and $q(\mathbf{x} \rightarrow \mathbf{x}')$

$$\pi_{t+1}(\mathbf{x}') = \sum_{\mathbf{x}} \pi_t(\mathbf{x}) q(\mathbf{x} \rightarrow \mathbf{x}')$$

Stationary distribution: $\pi_t = \pi_{t+1} = \pi$

$$\pi(\mathbf{x}') = \sum_{\mathbf{x}} \pi(\mathbf{x}) q(\mathbf{x} \rightarrow \mathbf{x}') \quad \text{for all } \mathbf{x}'$$

If π exists, it is unique (specific to $q(\mathbf{x} \rightarrow \mathbf{x}')$)

In equilibrium, expected “outflow” = expected “inflow”

Detailed Balance Equation

“Outflow” = “inflow” for each pair of states:

$$\pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') = \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x}) \quad \text{for all } \mathbf{x}, \mathbf{x}'$$

Detailed balance \Rightarrow stationarity:

$$\begin{aligned}\sum_{\mathbf{x}} \pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') &= \sum_{\mathbf{x}} \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x}) \\ &= \pi(\mathbf{x}') \sum_{\mathbf{x}} q(\mathbf{x}' \rightarrow \mathbf{x}) \\ &= \pi(\mathbf{x}')\end{aligned}$$

MCMC algorithms typically constructed by designing a transition probability q that is in detailed balance with desired π

Gibbs Satisfies Detailed Balance

Sample each variable in turn, given **all other variables**

Sampling X_i , let $\bar{\mathbf{x}}_i$ be all other nonevidence variables

Current values are x_i and $\bar{\mathbf{x}}_i$; \mathbf{e} is fixed

Transition probability is given by

$$q(\mathbf{x} \rightarrow \mathbf{x}') = q(x_i, \bar{\mathbf{x}}_i \rightarrow x'_i, \bar{\mathbf{x}}_i) = P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e})$$

This gives detailed balance with true posterior $P(\mathbf{x}|\mathbf{e})$:

$$\begin{aligned}\pi(\mathbf{x})q(\mathbf{x} \rightarrow \mathbf{x}') &= P(\mathbf{x}|\mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) = P(x_i, \bar{\mathbf{x}}_i | \mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) \\ &= P(x_i | \bar{\mathbf{x}}_i, \mathbf{e})P(\bar{\mathbf{x}}_i | \mathbf{e})P(x'_i | \bar{\mathbf{x}}_i, \mathbf{e}) \quad (\text{chain rule}) \\ &= P(x_i | \bar{\mathbf{x}}_i, \mathbf{e})P(x'_i, \bar{\mathbf{x}}_i | \mathbf{e}) \quad (\text{chain rule backwards}) \\ &= q(\mathbf{x}' \rightarrow \mathbf{x})\pi(\mathbf{x}') = \pi(\mathbf{x}')q(\mathbf{x}' \rightarrow \mathbf{x})\end{aligned}$$

Gibbs Sampling: Performance

Think of Gibbs sampling as

$$x^{l+1} \sim q(\cdot | x^l)$$

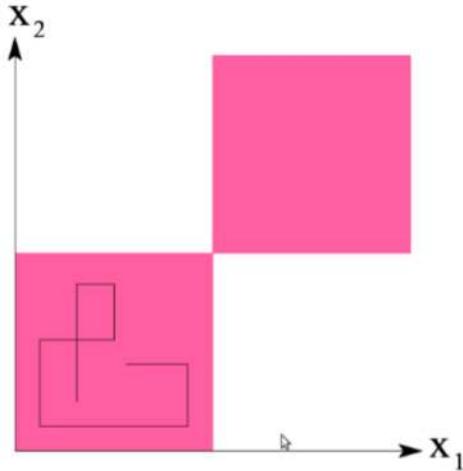
Problem: States are highly dependent (x^1, x^2, \dots)

Need a long time to run Gibbs sampling to *forget* the initial state, this is called **burn in** phase

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Gibbs_sampling

Gibbs Sampling: Performance



In this example the samples stay in the lower left quadrant

Some technical requirements to Gibbs sampling

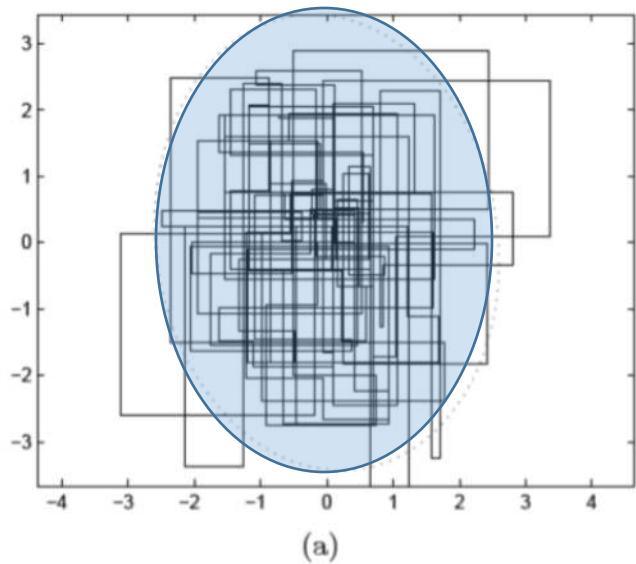
The Markov Chain $q(x^{l+1} | x^l)$ needs to be able to traverse the entire state-space (no matter where we start)

- ▶ This property is called **irreducible**
- ▶ Then $p(x)$ is the stationary distribution of $q(x' | x)$

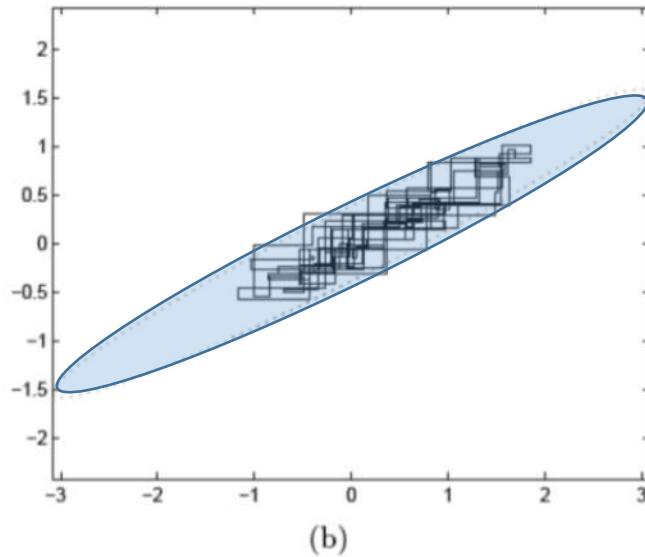
¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Gibbs_sampling

Gibbs Sampling: Performance



(a)



(b)

Gibbs sampling is more efficient if states are not correlated

- ▶ Left: Almost isotropic Gaussian
- ▶ Right: correlated Gaussian

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Gibbs_sampling

Metropolis-Hastings MCMC

- We will now mention one other MCMC method in passing.
 - Metropolis-Hasting (MH)
 - A special case is called Metropolis sampling.

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Metropolis%E2%80%93Hastings_algorithm

MH MCMC Algorithm

Slightly more general MCMC method when the proposal distribution is *not* symmetric

Sample x' and accept with probability

$$A(x', x) = \min \left(1, \frac{\tilde{q}(x | x') p^*(x')}{\tilde{q}(x' | x) p^*(x)} \right)$$

Note: when the proposal distribution is symmetric, Metropolis-Hastings reduces to standard Metropolis sampling

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Metropolis%E2%80%93Hastings_algorithm

MH MCMC Special Case: Metropolis Sampling

Special case of MCMC method (proposal distribution) with the following proposal distribution

- ▶ symmetric: $q(x' | x) = q(x | x')$

Sample x' and accept with probability

$$A(x', x) = \min \left(1, \frac{p^*(x')}{p^*(x)} \right) \in [0, 1]$$

- ▶ If new state x' is more probable always accept
- ▶ If new state is less probable accept with $\frac{p^*(x')}{p^*(x)}$

¹Reference: Bjoern Andres and Bernt Schiele, MPI (2016)

²Reference: https://en.wikipedia.org/wiki/Metropolis%E2%80%93Hastings_algorithm

Questions?

Summary

- Inference computations on joint distributions is a hard problem
- Graphical models help us do this in efficient ways
 - For tree models, this is linear time!
- We discussed two exact methods
 - Variable Elimination
 - Belief propagation
- We discussed one family of approximate methods
 - Based on **sampling** via Markov Chain Monte Carlo techniques

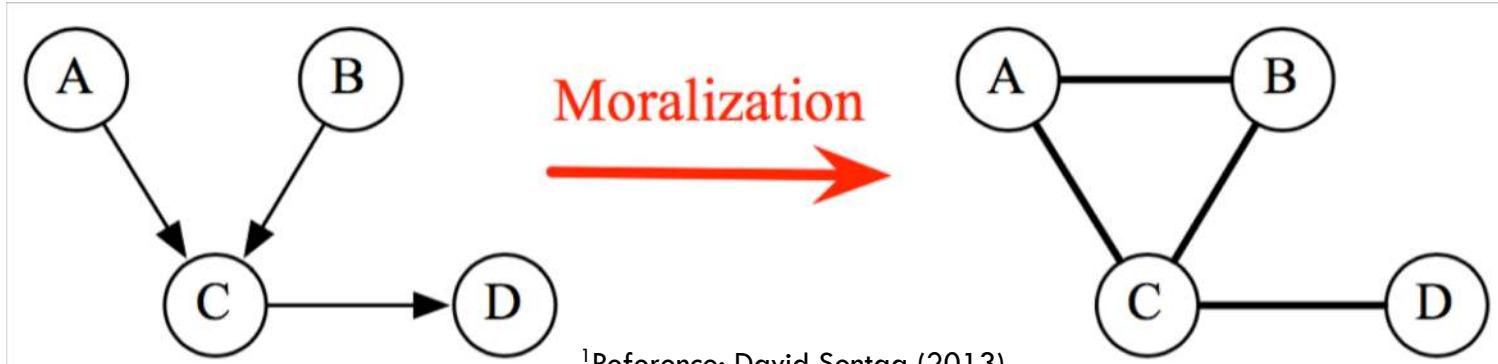
Appendix

Sample Exam Questions

- What is a factor graph? How is it related to DPGMs?
How is it related to UPGMs?
- What are the key steps of Belief propagation?
- What is the use of BP? Can it be used for inference over general factor graphs?
- How would one use sampling for inference?
- Why is Gibbs sampling a MCMC technique?
- Why does BP do better than variable elimination?

DPGMs and UPGMs

- Inference algorithms can typically run on both graphs
- For convenience, we will construct a UPGM from a DPGM and discuss inference on UPGM
- The construction is straightforward
 - For each factor in DPGM, call it a potential now
 - Moralize the DPGM and remove directions
 - (We lose some information in the graph)



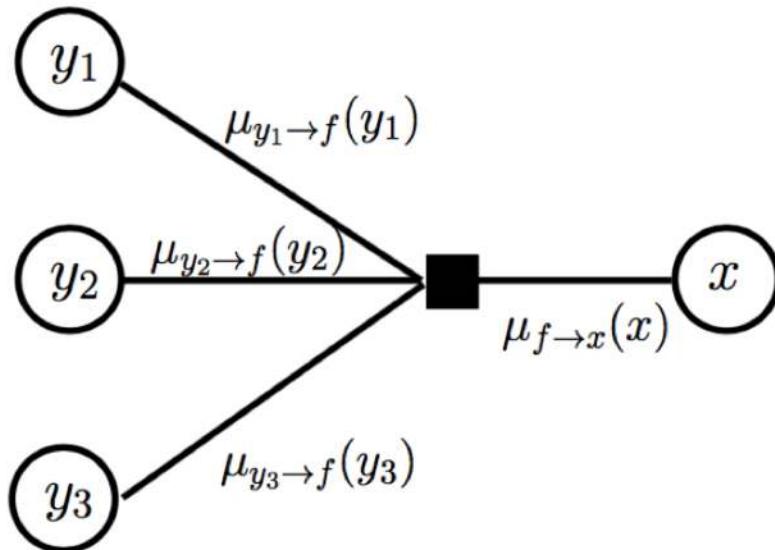
BP: Computing Maximal State

- BP variant can also solve for the maximal state $\bar{x}^* \in \operatorname{argmax}_{\bar{x} \in \mathfrak{X}} P(\bar{x})$
- This version is called **Max-Product Belief Propagation**
- Has three ingredients just as before
 - Initialization (same as before)
 - Variable to factor message (same as before)
 - Factor to variable message

BP: Computing Maximal State

- Factor to variable message is different from Sum-Product

$$\mu_{f \rightarrow x}(x) = \max_{y \in \mathcal{X}_f \setminus x} \phi_f(\mathcal{X}_f) \prod_{y \in \{\text{ne}(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$



- Additionally, we need to track values achieving maximums as well

BP: Computing Maximal State

- Maximal state of a variable is

$$x^* = \operatorname{argmax}_x \prod_{f \in \text{ne}(x)} \mu_{f \rightarrow x}(x)$$

